

Market Landscape & The Strategic Growth Opportunity Gap

WHY CLAUDE SKILLS

- Available on all paid plans: Pro, Max, Team, Enterprise
- Pre-configured within Claude.ai, no installation required
- Native access to full session output data, unreplicable by third parties
- High distribution. Near-zero activation.

CORE ISSUE

Claude delivers exceptional outputs but has no mechanism to learn from them. There is no persistence layer between a user and Claude. Preferences, corrections, and patterns are discarded after every session, forcing manual re-entry each time. The cost of this gap falls entirely on the user.

WHY THE CURRENT APPROACHES FAILS

- **Setup before value:** SKILL.md authoring required before any output benefit is shown
- **No contextual trigger:** Chat has never surfaced Skills at the moment of pain
- **Intent gap:** Every Skill creation flow asks users to describe preferences before seeing what great looks like
- **No pattern capture:** Correction behaviour goes unacknowledged by the product


PRODUCTS WITH SKILLS-LIKE PREFERENCE PERSISTENCE

- **Linear:** Remembers workflow preferences per user across sessions without any manual configuration
- **ChatGPT Memory:** Direct competitor shipping cross-session retention as a core feature. Table stakes, now
- **Zapier:** Captures repeated behaviour as a reusable rule. Trigger is action, not declared intent

HOW CLAUDE SKILLS COMPARES

FEATURE	CLAUDE SKILLS	OPENAI MEMORY	GEMINI	MS COPILOT	CURSOR
Remembers across sessions	Yes, via Skills	Yes, automatic	Yes, via Gems	Yes, via prompts	Yes, rules file
How it learns your preferences	You write them manually	It watches and learns	You configure upfront	You write a prompt	Learns from your work
Offered to you in the chat	Never	Yes	No	No	Yes
You can see and edit what it learned	Yes, fully visible	Partial	Yes	Yes	Yes
Asks before applying anything	Yes	No, applies automatically	Yes	Yes	Partial
What data it uses	Output structure only	Your full conversation	Config only	Config only	Codebase structure

THE SCALE OF OPPORTUNITY



At 20 to 30 million monthly active users, even a small improvement in session continuity translates into significant retention and conversion uplift.

Product Friction Audit: Deconstructing the Activation Breakdown

FUNCTIONAL CAPABILITIES OF CLAUDE SKILLS

- Saves tone, format, and structure preferences across sessions
- Surfaces reusable instruction sets for recurring, structured tasks
- Applies consistent rules to standardize output quality automatically
- Reduces repeated prompting by retaining formatting and style preferences

WHAT THE PRODUCT NEEDS TO DO

Shift the experience from manual re-prompting to automatic preference persistence by surfacing Skills at the moment of friction, when a user repeats context they have already given before, and the cost of starting over is most visible.

ENGAGEMENT AND TRUST BREAKDOWN


- **Effort friction:** SKILL.md setup demanded before any proof of value. Users who tried once abandoned and did not return
- **Cognitive load:** All preferences defined manually in technical syntax, no pattern capture, no guidance
- **Activation drop-off:** Paid users open Skills Settings reactively. No habit formed, no in-conversation trigger
- **Avoidance loop:** The pain the feature solves is recreated by the setup process itself

WHAT THE CURRENT UX OPTIMIZES

- **Fast lookup:** Quick access to saved Skills list for users who already built Skills
- **Compliance:** Supports manual instruction sets for technical users
- **Verification:** Confirms a Skill exists but does not help create one from a great output

It was not designed for users who, mid-conversation, receive a high-quality output and want to save it instantly for reuse.

THREE SCREENS WHERE CLAUDE FAILS ITS USERS



UX AUDIT-CURRENT CLAUDE SKILLS-FIVE FAILURE POINTS

FAILURE POINT	WHAT HAPPENS	USER IMPACT
Settings-first design	Skills requires navigating to Settings → Skills → New Skill before receiving any value	Investment precedes proof. Most users never reach this step
SKILL.md syntax barrier	Users must write structured instructions in a technical format before seeing any benefit	Technical friction filters out the users who need Skills most
No contextual trigger	Chat interface where the pain is felt, has never surfaced Skills to any user	The feature is invisible at the moment it would be most wanted
Passive discovery only	Users must scroll Settings without knowing Skills could automatically apply their preferences	Discovery happens outside the core workflow, preventing timely adoption
No pattern intelligence	No differentiation between a one-time great output and a recurring preference worth saving	Users cannot act on what the product cannot see

Research Framework: Validating Critical Activation Hypotheses

TARGET SEGMENT

Claude Pro or Max subscribers who use Claude for recurring professional tasks such as research synthesis, client writing, competitive analysis, structured reports at least 2× per week for the same category of task. Small as a proportion of all users but disproportionately responsible for retention, advocacy, and organic growth.

30 Survey respondents

5 Depth interviews

3 Recall walkthroughs

USER’S UNMET NEEDS

- Persistent preference memory across sessions without manual setup
- In-conversation Skill creation without leaving the chat interface
- Plain-English rules that are readable, editable, and inspectable
- Transparency into what Claude observed, surfaced in chat not buried in Settings
- Per-session toggle to disable a Skill without deleting it permanently

WHY THIS SEGMENT?

- These users return for the same task types repeatedly, making preference persistence immediately valuable
- These users bring the highest contextual input per session
- They work across multiple workstreams, each with distinct tone and format requirements Skills serves
- They engage in longer sessions. Longer sessions mean more context built up and more pain when that context disappears at session end
- Consistent output each time is a professional requirement for these users

THE THROUGH LINE CONNECTING ALL SIX

Users know what they want (H3), waste time not getting it (H4), build workarounds (H1), eventually abandon Claude for those tasks (H2), while a solution exists that Claude never surfaces (H6), that they would use if it gave them control (H5).

HYPOTHESES TESTED



	THE HYPOTHESIS	THE QUESTION BEING TESTED	KEY FINDING	VERDICT
H1	Users are already building workarounds on their own	If this is a real problem, users are already solving it themselves. Are they?	74% keep a separate doc with Claude instructions that gets pasted every session	Strongly Confirmed
H2	Inconsistent outputs cause users to quit Claude for specific tasks	Is the frustration bad enough to push users away entirely?	69% have stopped using Claude for a task that kept failing	Strongly Confirmed
H3	Users know what good looks like but cannot communicate it to Claude	Is the gap about knowledge or communication?	91% report having clear preferences. Only 13% feel Claude consistently understands them. A 78-point gap.	Strongly Confirmed
H4	The time lost to output corrections is quantifiable and significant	Can a concrete effort cost be attached to this problem?	85% need 2-3 attempts per task. 38% spend over 3 minutes correcting each output.	Confirmed
H5	Users will trust AI-learned preferences only with visible control over what is saved	How much control is needed before users are comfortable?	83% prefer a suggest-and-confirm model. Only 6% want automatic application.	Strongly Confirmed
H6	Skills has never been surfaced to a user who needed it, inside the product	Is this a discoverability gap or a product intelligence gap?	56% are aware of Skills. Every single one discovered it externally. Zero discovered it through Claude itself.	Partially Confirmed

Data Synthesis: Quantifying Frustration & Inconsistency

QUALITATIVE FINDINGS

- **The Quality Cliff:** Users describe receiving an exceptional output and then being unable to recreate it. The cause of the change is unknown to them. Self-blame follows, leading to progressively longer prompts. Output quality does not return
- **Scroll-Back Behaviour:** Users scroll chat history to locate the session where Claude produced the right output. This is a persistence problem, not a memory problem
- **Reflection Mismatch:** Users want to understand why output quality drifted across sessions, not simply how many corrections were made

THE PATTERN THAT KEEPS REPEATING

WHAT TRIGGERS IT	WHAT USERS DO	WHAT THEY THINK IS HAPPENING
Frustration after an output goes wrong	Scroll back through old chats to find a good session	They assume they prompted wrong, not that preferences were lost
Relief when it works temporarily	Start maintaining a document in their local device or Notion with prompt instructions	They think better prompting is the key to a good output
Confusion when identical tasks produce different results	Write longer and longer prompts	They accept inconsistency as how AI works
Anxiety before a high stake output	Spend 3 or more rounds correcting the prompt before sending	They never connect this to a missing product feature

Unprompted Upgrade Signal : 84% of interviewees stated they would upgrade from Free to Pro specifically for a feature that maintained consistent output quality. At approximately 20M monthly active users, even a 1% Free to Pro conversion lift from solving output consistency represents roughly \$3.4M in incremental ARR at the \$17/month Pro price point, making this the highest-leverage consumer monetization opportunity currently untested.

USER PERSONAS



Rujul K · Product Manager · Fintech · Mumbai · Claude Pro · 8 months

- Produces competitive analyses and investor ready strategy memos weekly
- Maintains a 400 word Claude briefing document in Notion that he pastes into every new session since 6 months
- Has not heard about Skills in Claude

"I got this perfect output once. The structure was exactly right, the tone was mine. I have been trying to get back to that ever since. I never have."



Vaibhav B · JS Developer · IT Consulting · Mumbai · Claude Pro · 4 months

- Uses Claude for code documentation and technical writing
 - Found Skills on YouTube. Tried it once. The blank editor stopped him.
 - Spends 15 minutes re-explaining his style every single session
- "The blank box with no guidance felt like writing code from scratch. I gave up and went back to re-prompting."

USER'S RECALL BASED WALKTHROUGH

- I use Claude almost every day for work
- I always end up writing the same instructions, same format, same tone, same structure
- I found something called Skills in the settings once. I clicked on it
- There was a blank box. I did not know what to write. I closed it
- The next day I opened a new chat and typed my instructions again from scratch
- A colleague told me Skills was supposed to remember my preferences automatically
- I went back and tried again. The blank editor still gave me nothing to start with
- I gave up and went back to pasting my old prompt at the start of every session
- I know the feature exists. I just never felt like it was built for me
- If Claude had just asked me two questions after a good output, I would have set it up that day

Problem Framing Canvas

WHAT IS THE TRUE PROBLEM?

Recurring structured tasks demand consistent output quality, yet users re-enter the same format, tone, and structure preferences every session. Claude has a native preference layer built to solve exactly this, but it has never been triggered mid-conversation. The capability exists within the product. The mechanism to surface it at the moment it would matter most does not.

WHO ARE THE CUSTOMERS FACING THE PROBLEM?

Recurring Task Power Users

- **Subscription Tier:** Hold paid Pro or Max subscriptions and use Claude daily
- **Behavior:** Perform the same category of structured tasks repeatedly at least 2x per week for the same category of task
- **Workarounds:** They have already built external manual memory systems, such as Notion prompt libraries, saved chats, and copy-paste templates
- **Status:** They are motivated and pattern-blind, frustrated that the tool does not learn their specific standards

HOW DO WE KNOW IT IS A REAL PROBLEM?

- **Qualitative:** Quality cliff described across all 5 interviews. Users cannot recreate their best outputs. They scroll back through chat history. They blame themselves. They accept inconsistency as normal
- **Quantitative:** 69% abandonment. 74% workarounds. 38% spend 3+ minutes per correction. 85% need 2–3 attempts per task. The cost is measurable and large

WHAT IS THE VALUE GENERATED FOR THE CUSTOMERS?

- Output matches their standard on the first attempt, every time
- The correction loop is permanently broken. No re-prompting, no pasted instructions
- Every preference is readable and editable before the Skill goes live
- Capture happens at peak satisfaction, right after a great output
- Users shift from managing Claude to relying on it

WHAT IS THE VALUE GENERATED FOR THE BUSINESS?

- Skills-activated users churn less. Each Skill creates a personalized switching cost
- Free users who experience Skills-quality output are more willing to convert to paid plans
- Correction and preference signals build a compounding dataset that grows more defensible over time
- Confirm-first design is a trust advantage over OpenAI memory which applies silently
- Skills is already live. Discovery improvements convert existing capability into retention

WHY SOLVE THIS PROBLEM NOW?

- **Competitive Urgency:** OpenAI is already shipping "Memory" to GPT-4o. As competitors add persistence, the perception gap regarding Claude's capability grows
- **Untapped Revenue:** 84% of free users indicated they would consider upgrading if Claude remembered their preferences-this is a massive, unexploited conversion lever
- **Time-Sensitivity:** The first AI to build reliable preference learning at scale creates compounding advantages that are difficult for late-movers to disrupt

Proposed Solutions

Solution 1: Skill from Output (Selected)



WHAT IT DOES

When Claude delivers a structurally strong output, or detects a second correction on the same task type, a Save as Skill chip appears inline alongside the standard thumbs-up and copy icons. Tapping opens a plain-English summary of what Claude learned from that output. Behind the surface, Claude Haiku analyses the response and extracts format rules; Claude Sonnet compiles them into a complete Skill. The user reviews, edits if needed, and confirms. The great output becomes the instruction.

WHY IT MIGHT WORK

- Starts from evidence, not intent; encodes what great actually looks like (H3)
- Surfaces at the highest-intent moment in any session, immediately post-great output (H6)
- Suggest-and-confirm design directly matches H5's 83% preference finding
- Eliminates the quality cliff by making the exceptional output the new baseline (H2)

HYPOTHESES ADDRESSED

H2 Abandonment, H3 Communication gap, H5 Trust model and H6 Discovery

Solution 2: Onboarding Interview

WHAT IT DOES

Claude conducts a 3-question structured conversation when a new Pro user signs up: What do you use me for? What does a great output look like? What should I never do? From the answers, Claude generates a draft Skill in plain English, shows a preview, and asks for confirmation before activating.

WHY IT MIGHT WORK

- Creates a Skill on day one, before any frustration accumulates
- Requires no great output to exist yet; works at account creation
- Confirm-first design aligns with H5 trust requirement
- Surfaces Skills proactively; closes H6 for new users

HYPOTHESES ADDRESSED

H3 Communication gap and H6 Discovery

WHY IT WAS DEPRIORITISED

This is the right solution if the primary problem is new user activation. It was deprioritized because H3 data shows a 78-point gap between stated and implicit preferences. The Skill it produces is built on what users think they want, not what great actually looks like a quality ceiling that limits long-term retention value.

Solution 3: Skill Try-out Sandbox

WHAT IT DOES

Before installing any Skill, the user pastes a real piece of their own work into a sandboxed preview. The Skill runs against it and the user sees exactly what it produces on their actual input before committing. Installation occurs only after seeing real output.

WHY IT MIGHT WORK

- Directly addresses H5 trust concern: applying wrong preferences to new context
- Converts uncertain discovery into confident installation
- No commitment required before proof of value
- Closes the gap between Skills discovery and Skills adoption

HYPOTHESES ADDRESSED

H5 Trust model and H6 Discovery

WHY IT WAS DEPRIORITISED

This is the right solution if the primary problem is Skill adoption from a shared library. It directly addresses H5's trust concern. It was deprioritized because it is downstream of creation; personalized Skills do not exist yet for most users, so there is nothing to try out. This becomes the right V1.1 feature once Skill from Output has seeded the library.

Solution Prioritization and Design

PRIORTIZATION FRAMEWORK: RICE

SOLUTION	REACH (1-10)	IMPACT (1-3)	CONFIDENCE	EFFORT (WEEKS)	RICE SCORE	HYPOTHESES
✓ Skill from Output	9	3.0	85%	6	38.25	H2, H3, H5, H6
Skill Tryout Sandbox	6	2.0	70%	5	16.80	H5, H6 only
Onboarding Interview	5	1.5	55%	4	10.31	H3, H6 only


Score = (Reach x Impact x Confidence) / Effort. Reach rated 1-10 based on proportion of eligible Pro users exposed per quarter. Impact: 3 = solves the core problem, 2 = meaningfully helpful, 1 = minor improvement. Confidence based on research alignment across all six hypotheses.

WHY THIS SOLUTION WINS ON EVERY DIMENSION

- Reach is highest because the chip appears for every Pro or Max user who gets a qualifying output. No opt-in required. It reaches the full eligible base automatically
- Impact is must-have because it is the only solution that solves both the quality cliff (H2) and the discovery gap (H6) at the same time. No other option does both
- Confidence is highest because three independent data points all point to the same design. H3 says capture from evidence not intent. H5 says confirm before applying. H6 says surface it in the chat
- Effort is manageable because the pattern detection runs on Claude's existing session data. No new infrastructure and no new data pipeline required


"Every other approach asks users what they want before they have seen it. This one waits until they have seen it, then offers to save it."

USER FLOW DIAGRAM




KEY
Discovery (H6) Trust / Control (H5, H3) Retention (H2)
--- Alt / dismissed path

SYSTEM ARCHITECTURE





Output structure only · On-session · User-controlled



Wireframing and Product Flow





Screen 1: Great Output Claude detects a high-quality output from a recurring task type and surfaces 'Save as Skill?' with no user action needed.




Screen 2: Analyzer Running Claude Haiku detects 6 preferences from the output automatically. Zero effort from the user.




Screen 3: SKILL.md Generated Rules extracted from the output, not typed by hand. The user sees exactly what was captured.




Screen 4: SKILL.md Preview + Edit Full transparency before commitment. Read and edit every rule before the Skill goes live.





Screen 5: The user can try the Skill on a sample before committing. See it work before saving it.



Screen 6: One keyword set. Claude will activate this Skill whenever that phrase appears.






Screen 7: The user types a similar message and Claude detects it matches the saved Skill pattern.



Screen 8: Claude silently activates the Skill in the background and applies all saved rules to the response.

Screen 9: The user receives a Skill-formatted output and chooses: 'Looks good' or 'Not quite.'



Screen 10: Selecting 'Not quite' surfaces options to adjust tone, format, length or structure so the Skill improves.



Measurement: Growth Indicators & Guardrail Metrics

NORTH STAR METRIC

Skill Activation Rate among Pro/Max users

% of eligible Pro/Max sessions resulting in a confirmed Skill within 30 days of being shown the chip.

Skill Activation Rate = (Confirmed Skill) / (Users shown the chip) x 100

Target : 25% of exposed users in 30 days

*Chip refers to the Save as Skill banner visible to users

LEADING INDICATORS: 0 to 30 DAYS

METRICS	WHAT IT INDICATES	TARGET
Skill chip click through rate	Whether the chip feels relevant and worth exploring	>50% of exposed users
Rule confirmation rate	Whether Claude is extracting the right preferences from the output	>60% confirm & save
Context edit rate	Whether users feel genuine ownership, not passive	>40% edit at least 1 rule
Skill Day 7 retention	Whether it actually improved quality enough to keep	>70% use Skill in week 2

GUARDRAIL METRICS

- If most users see the chip and dismiss it without tapping, the chip is not landing right.
Skill Disable Rate <10% - Prevents perception of over-automation
- If users confirm Skills but then disable them within a day, the rules are not accurate enough
Rule Rejection Rate <30% - If too high, compiler is generating wrong rules
- If users feel surveilled or uncomfortable, there's failure on the privacy framing regardless of other metrics
Privacy Concern Reports <1% - Ensures session-level analysis is trusted

LAGGING INDICATORS: 30 to 90 DAYS

METRICS	WHAT IT INDICATES	TARGET
Pro/Max Plan Churn Delta	Encoded preferences make switching structurally harder over time.	8–12% reduction
Skill 30-Day Retention	Measures whether Skill output quality sustains engagement.	>55%
Free to Pro/Max Conversion Uplift	Locked chip exposure converts feature desire into paid plan upgrades.	+8 to 12% lift vs non-exposed free users
Skills WAU Penetration	Measures sustained platform-level adoption	25–30% of eligible Pro/Max users

MONETIZATION

- **Free to paid conversion:** Free users see the locked chip at peak output quality. The product creates the upgrade moment without any marketing spend
- **Pro to Team upsell:** In-chat Skill creation gives users something worth sharing. Skill sharing requires Team plan. The activation flow creates the upsell trigger
- **Preference editing is a paid-only capability:** Full control over saved rules, tone, and format is exclusive to paid plans. Surfacing this in conversation makes the paid plan tangibly more valuable than the free tier

DISTRIBUTION

- **In-conversation placement** at peak output quality eliminates the five-tap Settings discovery barrier entirely.
- **Automatic pattern detection** surfaces the chip without requiring any user-initiated action or prior awareness of Skills.
- **Every active Pro and Max session** is a native distribution touchpoint, requiring no external channel or campaign.

Full Rollout Cycles and Privacy Protections

GO-TO-MARKET

PHASE 1: WEEKS 1-4

Controlled Pilot (Trust and Accuracy Validation)

Audience: Pro users with 5+ recurring task types in the last 30 days

Rollout: Opt-in chip in Claude chat for 4 weeks. Surface chip on top-quality outputs only; conservative confidence threshold

Track: Confirmation rate, disable rate, correction loop reduction

Kill threshold: Chip rejection rate exceeds 30%

Success criteria: SAR above 20% · Disable rate below 10%

PHASE 2: WEEKS 5-10

Targeted Expansion (Relevance Scaling)

Audience: Pro users with detectable recurring patterns, 2+ same task type per week.

A/B test chip placement: Inline beneath output vs. message action bar

Micro-survey: Post-Skill-creation for qualitative feedback

Success criteria: Skills WAU increase of 20% vs baseline · 15% Skill creation rate among chip-impression users

PHASE 3: WEEKS 11-16

Full Rollout and Upgrade Trigger

Audience: All Pro, Max, Team users plus Free tier with locked state

Chip visible in Free tier with locked state: Converts feature moment into upgrade trigger

Introduce Skill Gifting as V1.1: Share a Skill via link for peer-driven adoption

Success criteria: 5% Free-to-Pro conversion from chip exposure

ETHICAL GUIDELINES

#	PRINCIPLE	WHAT THIS MEANS IN PRACTICE
1	Explicit Informed Consent	Before the chip ever surfaces, users understand that it analyses output structure to suggest preferences
2	On-session processing	All pattern detection happens within the current session. No content leaves. No cross-session data without explicit confirmation
3	No output monetisation	Saved preferences are never used for ads, targeting, or profiling. The feature exists solely to improve output quality for that user
4	Anonymized research data	Everything from the survey and interviews is anonymized. No names. No identifiable information stored
5	Emotional safety by design	The system surfaces suggestions, never verdicts. Users can confirm, edit, dismiss, or delete any Skill at any time
6	Right to opt out	The chip can be turned off in one tap. Existing Skills are exportable. No lock-in of any kind

ASSOCIATED RISKS AND MITIGATIONS

RISK	MITIGATION
The chip feels intrusive	Explicit 'here is what data was used' module. On-session processing only. Clear opt-out always visible
The extracted rules are wrong	Confidence threshold before the chip even appears. Full edit rights. Manual fallback if needed
Users confirm but quality does not improve	Track correction rate in Skill-active users. If it does not drop, the rules are not the right rules
The Skill becomes stale months later	Monitor edit frequency after activation. If users keep changing outputs even with a Skill active, the Skill needs updating.